

0 | 0 0 | 0 0 0 | 0

Pandas Tutorial 2: Aggregation and Grouping

Written by Tomi Mester on July 23, 2018

Let's continue with the pandas tutorial series. This is the second episode, where I'll introduce aggregation (such as min, max, sum, count, etc.) and grouping. Both are very commonly used methods in analytics and data science projects – so make sure you go through every detail in this article!

Note 1: this is a hands-on tutorial, so I recommend doing the coding part with me!

Before we start

If you haven't done so yet, I recommend going through these articles first:

1. [How to install Python, R, SQL and bash to practice data science](#)
2. [Python for Data Science – Basics #1 – Variables and basic operations](#)
3. [Python Import Statement and the Most Important Built-in Modules](#)
4. [Top 5 Python Libraries and Packages for Data Scientists](#)
5. [Pandas Tutorial 1: Pandas Basics \(Reading Data Files, DataFrames, Data Selection\)](#)

Aggregation is the process of turning the values of a dataset (or a subset of it)

<u>animal</u>	<u>water need</u>
zebra	100
lion	350
elephant	670
kangaroo	200

...then a simple aggregation method is to calculate the summary of the water_needs, which is $100 + 350 + 670 + 200 = 1320$. Or a different aggregation method would be to count the number of the animals, which is 4. So the theory is not too complicated. Let's see the rest in practice...

Data aggregation – in practice

Where did we leave off [last time](#)? We opened a Jupyter notebook, imported pandas and numpy and loaded two datasets: `zoo.csv` and `article_reads`. We will continue from here – so if you haven't done the "[pandas tutorial – episode 1](#)", it's time to go through it!

Okay!

Let's start with our `zoo` dataset! (If you want to download it again, you can find it at [this link](#).) We have loaded it by using:

```
pd.read_csv('zoo.csv', delimiter = ',')
```

```
In [1]: import numpy as np  
import pandas as pd
```



```
In [4]: pd.read_csv('zoo.csv', delimiter = ',')
```

```
Out[4]:   animal  uniq_id  water_need  
0 elephant    1001      500  
1 elephant    1002      600  
2 elephant    1003      550  
3 tiger       1004      300  
4 tiger       1005      320  
5 tiger       1006      330  
6 tiger       1007      290  
7 tiger       1008      310  
8 zebra        1009      200  
9 zebra        1010      220  
10 zebra       1011      240
```

Let's store this dataframe into a variable called `zoo`.

```
zoo = pd.read_csv('zoo.csv', delimiter = ',')
```

```
In [1]: import numpy as np  
import pandas as pd
```



```
In [5]: zoo = pd.read_csv('zoo.csv', delimiter = ',')
```

Okay, let's do five things with this data:

1. Let's count the number of rows (the number of animals) in `zoo` !
2. Let's calculate the total `water_need` of the animals!
3. Let's find out which is the smallest `water_need` value!
4. And then the greatest `water_need` value!
5. And eventually the average `water_need` !

Pandas Data Aggregation #1: `.count()`

Counting the number of the animals is as easy as applying a `count` function on the `zoo` dataframe:

```
zoo.count()
```

```
In [14]: zoo.count()
Out[14]: animal      22
          uniq_id    22
          water_need  22
          dtype: int64
```

Oh, hey, what are all these lines? Actually, the `.count()` function counts the number of values in *each column*. In the case of the `zoo` dataset, there were 3 columns, and each of them had 22 values in it.

If you want to make your output clearer, you can select the `animal` column first by using one of the selection operators from the previous article:

```
zoo[['animal']].count()
```

```
In [17]: zoo[['animal']].count()
Out[17]: animal    22
          dtype: int64
```

Or in this particular case, the result could be even nicer if you use this syntax:

```
zoo.animal.count()
```

This also selects only one column, but it turns our pandas dataframe object into a pandas series object. (*Which means that the output format is slightly different.*)

```
In [15]: zoo.animal.count()
Out[15]: 22
```

Pandas Data Aggregation #2: `.sum()`

Following the same logic, you can easily sum the values in the `water_need` column by typing:

```
zoo.water_need.sum()
```

```
In [19]: zoo.water_need.sum()
```

```
Out[19]: 7650
```

Just out of curiosity, let's run our sum function on all columns, as well:

```
zoo.sum()
```

```
In [20]: zoo.sum()
```

```
Out[20]: animal      elephantelephantelephantigertigertigerti...  
         uniq_id           22253  
         water_need          7650  
         dtype: object
```

Note: I love how `.sum()` turns the words of the `animal` column into one string of animal names. (By the way, it's very much in line with the logic of Python.)

Pandas Data Aggregation #3 and #4: `.min()` and `.max()`

What's the smallest value in the `water_need` column? I bet you have figured it out already:

```
zoo.water_need.min()
```

```
In [21]: zoo.water_need.min()
```

```
Out[21]: 80
```

And the max value is pretty similar:

```
zoo.water_need.max()
```

```
In [22]: zoo.water_need.max()
```

```
Out[22]: 600
```

Pandas Data aggregation #5 and #6: `.mean()` and `.median()`

Eventually, let's calculate **statistical averages**, like mean and median:

```
zoo.water_need.mean()
```

```
In [25]: zoo.water_need.mean()  
Out[25]: 347.72727272727275
```

```
zoo.water_need.median()
```

```
In [26]: zoo.water_need.median()  
Out[26]: 325.0
```

Okay, this was easy. Much, much easier than the aggregation methods of SQL.
But let's spice this up with a little bit of grouping!

Grouping in pandas

As a Data Analyst or Scientist you will probably do segmentations all the time. For instance, it's nice to know the mean `water_need` of all animals (we have just learned that it's `347.72`). But very often it's much more actionable to break this number down – let's say – by animal types. With that, we can compare the species to each other – or we can find outliers.

Here's a simplified visual that shows how pandas performs “segmentation” (grouping and aggregation) based on the column values!

animal	water_need	
zebra	10	
zebra	15	
lion	100	
elephant	320	
zebra	20	
lion	120	
lion	140	
zebra	15	

zebra → mean: 15

lion → mean: 120

elephant → mean: 320

Pandas .groupby in action

Let's do the above presented grouping and aggregation for real, on our `zoo` DataFrame!

We have to fit in a `groupby` keyword between our `zoo` variable and our `.mean()` function:

```
zoo.groupby('animal').mean()
```

```
In [32]: zoo.groupby('animal').mean()
Out[32]:
      uniq_id  water_need
      animal
      elephant  1002.0  550.000000
      kangaroo  1021.0  416.666667
      lion     1017.5  477.500000
      tiger    1006.0  310.000000
      zebra    1012.0  184.285714
```

Just as before, pandas automatically runs the `.mean()` calculation for all remaining columns (the `animal` column obviously disappeared, since that was

the column we grouped by). You can either ignore the `uniq_id` column, or you can remove it afterwards by using one of these syntaxes:

`zoo.groupby('animal').mean()[['water_need']]` (This returns a DataFrame object.)

`zoo.groupby('animal').mean().water_need` (This returns a Series object.)

```
In [39]: zoo.groupby('animal').mean()[['water_need']]
```

```
Out[39]: water_need
```

animal	
elephant	550.000000
kangaroo	416.666667
lion	477.500000
tiger	310.000000
zebra	184.285714

Obviously, you can change the aggregation method from `.mean()` to anything we learned above!

Okay! Now you know everything, you have to know!
It's time to...

Test yourself #1

Let's get back to our `article_read` dataset.

(Note: Remember, this dataset holds the data of a travel blog. If you don't have the data yet, you can download it from [here](#). Or you can go through the whole download, open, store process step by step by reading the previous episode of this [pandas tutorial](#).)

In [46]: article_read

Out[46]:

	my_datetime	event	country	user_id	source	topic
0	2018-01-01 00:01:01	read	country_7	2458151261	SEO	North America
1	2018-01-01 00:03:20	read	country_7	2458151262	SEO	South America
2	2018-01-01 00:04:01	read	country_7	2458151263	AdWords	Africa
3	2018-01-01 00:04:02	read	country_7	2458151264	AdWords	Europe
4	2018-01-01 00:05:03	read	country_8	2458151265	Reddit	North America
5	2018-01-01 00:05:42	read	country_6	2458151266	Reddit	North America
6	2018-01-01 00:06:06	read	country_2	2458151267	Reddit	Europe
7	2018-01-01 00:06:15	read	country_6	2458151268	AdWords	Europe
8	2018-01-01 00:07:21	read	country_7	2458151269	AdWords	North America
9	2018-01-01 00:07:29	read	country_5	2458151270	Reddit	North America
10	2018-01-01 00:07:57	read	country_5	2458151271	AdWords	Asia

If you have everything set, here's my first assignment:

What's the most frequent source in the article_read dataframe?

.

.

.

And the solution is: **Reddit!**

How did I get it? Use this code:

```
article_read.groupby('source').count()
```

Take the article_read dataset, create segments by the values of the source column (groupby('source')), and eventually count the values by sources (.count()).

In [47]: article_read.groupby('source').count()

Out[47]:

source	my_datetime	event	country	user_id	topic
AdWords	500	500	500	500	500
Reddit	949	949	949	949	949
SEO	346	346	346	346	346

You can – optionally – remove the unnecessary columns and keep the user_id column only:

```
article_read.groupby('source').count()[['user_id']]
```

Test yourself #2

Here's another, slightly more complex challenge:

For the users of `country_2`, what was the most frequent topic and source combination? Or in other words: which topic, from which source, brought the most views from `country_2`?

.

.

.

The result is: the combination of Reddit (source) and Asia (topic), with 139 reads!

And the Python code to get this results is:

```
article_read[article_read.country ==  
'country_2'].groupby(['source', 'topic']).count()
```

In [58]: article_read[article_read.country == 'country_2'].groupby(['source', 'topic']).count()					
Out[58]:					
		my_datetime	event	country	user_id
AdWords	Africa	3	3	3	3
	Asia	31	31	31	31
	Australia	6	6	6	6
	Europe	46	46	46	46
	North America	11	11	11	11
	South America	14	14	14	14
Reddit	Africa	24	24	24	24
	Asia	139	139	139	139
	Australia	18	18	18	18
	Europe	29	29	29	29
	North America	27	27	27	27
	South America	26	26	26	26
SEO	Africa	7	7	7	7
	Asia	9	9	9	9
	Australia	10	10	10	10
	Europe	4	4	4	4
	North America	42	42	42	42
	South America	16	16	16	16

Here's a brief explanation:

First, we filtered for the users of `country_2` (`article_read[article_read.country == 'country_2']`). Then on this subset, we applied a `groupby` pandas method... Oh, did I mention that you can group by multiple columns? Now you know that! 😊 (Syntax-wise, watch out for

one thing: you have to put the name of the columns into a list. That's why the bracket frames go between the parentheses.) (That was the `groupby(['source', 'topic'])` part.)

And as per usual: the `count()` function is the last piece of the puzzle.

Conclusion

This was the second episode of my pandas tutorial series. I hope now you see that aggregation and grouping is really easy and straightforward in pandas... and believe me, you will use them a lot!

Note: If you have used SQL before, I encourage you to take a break and compare the pandas and the SQL methods of aggregation. With that you will understand more about the key differences between the two languages!

In the next article, I'll show you the four most commonly used "data wrangling" methods: `merge`, `sort`, `reset_index` and `fillna`. Stay with me: [Pandas Tutorial, Episode 3!](#)

- If you want to learn more about how to become a data scientist, take my 50-minute video course: [How to Become a Data Scientist](#). (It's free!)
- Also check out my 6-week online course: [The Junior Data Scientist's First Month video course](#).
- Or simply subscribe to my [Newsletter](#).

Cheers,

Tomi Mester

Help your friends see this article:



[← PREVIOUS POST](#)

[NEXT POST →](#)

2 Comments



Nigel Ogbuchi

FEBRUARY 5, 2019

Great post. learnt a few tricks from it. Thank you.

[REPLY](#)



Tomi Mester

MARCH 6, 2019

Cool! Glad it helped! :)

And thanks for the comment!

[REPLY](#)

Leave a Reply

[COMMENT](#)

NAME *

EMAIL *

WEBSITE

Post Comment